ADI Anvil

Founder brief | 2 November 2026

Business-context benchmarking that turns model failures into evidence for improvement.

The problem

Frontier labs and enterprises need to know whether a model can do real business work, where it fails and what to improve. A benchmark score alone cannot establish whether a decision respected a bank's KYC policy, a campaign's budget or the evidence available at that moment. Generic tasks often omit the rules and context that make those decisions consequential.

What we are building

Anvil is an evaluation and diagnostics platform for testing models against versioned business tasks. It connects performance to the underlying evidence, constraints, rubrics and observable actions. Labs and enterprises can investigate gaps, compare systems under the same conditions and identify task families and examples that could support their own fine-tuning or training. Anvil does not currently train model weights or guarantee improvement.

Scale, trust and business context

Scale: Anvil is designed to reuse task collections, execution environments and grading rules across models and repeated evaluations. Our next step is to grow Marketing, Banking and Commerce taskbanks toward 200-500 expert-authored samples each, then expand execution capacity. These are targets, not today's inventory or demonstrated throughput.

Trust: Exact task, rubric, model and environment versions make results traceable. Access controls, separated hidden references, budget limits and retained failure evidence support integrity. Planned compliance checks will link task rules to applicable laws, jurisdictions and specialist review, flagging unresolved requirements. They will not constitute automatic legal certification.

Business context: Praxis supplies the expert-authored tasks behind the tests. Our intended advantage is the chain from a verified practitioner's judgment to a reviewable model result: who defined the rule, why it matters and how the model violated it. A leaderboard score alone cannot verify that chain of business evidence.

Progress and next milestones

Synthetic Banking KYC/AML, fraud and Marketing tasks run through the platform. Astra, Claude and Gemini have been exercised; the latest bounded Gemini diagnostic completed successfully. The Benchmarks library and comparison dashboard are explicitly labeled UI mockups, with illustrative values and full functionality deferred.

Over the next 30 days, priorities are remaining product functionality, SME recruitment and the planned Anvil-to-Praxis review handoff with decisions logged back into Anvil. Beyond December 2026, we plan integrations with Prammana and Thinking Machines to support the broader evaluation-to-improvement workflow. These are roadmap intentions, not implemented integrations or confirmed partnerships.

Together: Praxis makes business expertise explicit. Anvil makes model behavior against that expertise inspectable.