ADI Praxis

Founder brief | 2 November 2026

Expert task authoring that makes real-world business judgment usable for AI evaluation.

The problem

AI tests are only as trustworthy as the tasks and grading rules behind them. A plausible question can miss the operating constraints, exceptions and evidence that an experienced practitioner uses. Without knowing who authored a task and who verified its rules, a large dataset can scale uncertainty instead of useful judgment.

What we are building

Praxis is a domain-expert task-authoring platform. It structures business work into tasks, evidence, rules, rubrics, constraints and expected deliverables that Anvil can test. Our task-sourcing work draws on conversations with CMOs, SEO analysts, campaign managers, and KYC and compliance managers at banks across geographies. The goal is to capture how real decisions are made, rather than rely only on generic or model-generated questions.

Scale, trust and business context

Scale: A repeatable authoring and review process lets many experts contribute while preserving consistent task structure. We aim to build Marketing, Banking and Commerce taskbanks of 200-500 independently expert-authored samples per domain, using Banking KYC/AML and fraud as reference use cases. These are recruitment and publication targets, not completed collections.

Trust and integrity: Our production verification standard requires rigorous SME onboarding: identity, relevant experience, domain fit, conflicts and contribution rights. Each task should receive documented expert verification before publication. We intend to involve senior specialists, including retired Federal Reserve or RBI officials where engaged, and professors or guides suited to the task. Participation is not regulatory approval or institutional endorsement. Production-grade verification and the full review process remain milestones to complete.

Business context and compliance: Experts author the rules, rubrics and constraints that distinguish a valid answer from an unsafe or impractical one. Reviews should identify the applicable jurisdiction, policy or legal source, exceptions and review date. Version history, review decisions and permission boundaries create a trail that can be examined and updated. This supports compliance review; it does not guarantee compliance with every law.

Why this can be different

Our intended advantage combines access to practitioners with a verifiable authoring process. A collection of questions alone cannot establish the author's qualifications, the rule's business basis or the review behind a grading decision. Praxis is designed to preserve those connections as contributions scale. Proposed recruiting channels include the Bodhium network for Marketing and our Banking network; university and former-regulator participation must be confirmed individually.

Progress and next milestones

The current platform supports structured authoring, review and versioned exports for Anvil, with synthetic reference tasks. Over the next 30 days, we will focus on remaining functionality, expert recruitment, verification and the planned Anvil run-to-SME-review-to-Anvil feedback loop. Beyond December 2026, the broader product roadmap includes proposed Prammana and Thinking Machines integrations; scope and partner participation are unconfirmed.

Together: Praxis establishes the business test. Anvil shows where the model meets it, fails it and needs further investigation.